

Conjectural Tensor Rank Bounds via Non-Commutative Algebras and Nilpotent Extensions

Xavier Callens

Socrate AI Lab, Paris, France

xavier@socrateai.org

June 2026

Abstract

We propose a framework for bounding tensor rank via non-commutative algebras augmented with nilpotent elements. We define a 4-dimensional algebra $\mathcal{A}$ over $\mathbb{R}$ with basis $\{1, \mathbf{i}, \mathbf{j}, \varepsilon\}$, where $\varepsilon^2 = 0$, and prove in the Lean 4 proof assistant that the commutator of any two elements has vanishing trace and is concentrated entirely in the ε -channel. Building on this verified algebraic property, we *conjecture* that the border rank of the $\langle N, N, N \rangle$ matrix multiplication tensor satisfies $\underline{R}(\langle N, N, N \rangle) \leq 4N^2(\lceil \log_2 N \rceil + 1)$, which via Schönhage’s τ -theorem would imply $\omega = 2$. We also propose an axiom decomposition of the lace expansion for self-avoiding walks on $\mathbb{Z}^3$. All unproven claims are stated explicitly as conjectures. We present seven concrete open problems.

MSC 2020: 68Q17, 15A69, 16S50, 82B41.

Keywords: matrix multiplication exponent, tensor rank, border rank, non-commutative algebra, nilpotent extension, self-avoiding walks, open problems.

1 Introduction

The exponent of matrix multiplication, denoted ω , is the infimum over all real numbers α such that two $N \times N$ matrices can be multiplied using $O(N^\alpha)$ arithmetic operations. Determining the exact value of ω is among the most important open problems in algebraic complexity theory.

The trivial algorithm gives $\omega \leq 3$. Strassen’s celebrated 1969 algorithm [19] computes 2×2 matrix products using 7 multiplications rather than 8, establishing $\omega \leq \log_2 7 \approx 2.807$. Subsequent work by Pan [16], Bini et al. [3], Schönhage [17], and Coppersmith–Winograd [10] introduced the *laser method* and the τ -theorem, progressively lowering the upper bound. The current best result is

$$\omega < 2.371552, \tag{1}$$

due to Duan, Wu, and Zhou [11], building on the breakthrough of Alman and Williams [1]. The information-theoretic lower bound $\omega \geq 2$ is elementary: one must at minimum *read* the $2N^2$ input entries. Whether $\omega = 2$ remains wide open.

Very recently, DeepMind’s AlphaEvolve system [2] discovered a tensor decomposition of rank 48 for the $\langle 4, 4, 4 \rangle$ matrix multiplication tensor, improving the previous best of 49 (from Strassen’s recursive construction). This result is *verified*: the rank-48 decomposition has been checked by explicit symbolic computation.

Contributions and limitations. In this paper, we:

1. Define a 4-dimensional non-commutative algebra with a nilpotent generator and *prove* (via Lean 4) that commutator traces vanish (Section 3).
2. *Conjecture* a border rank bound $\underline{R}(\langle N, N, N \rangle) \leq 4N^2(\lceil \log_2 N \rceil + 1)$ and discuss its implications via Schönhage's τ -theorem (Section 2).
3. Present a partial tensor decomposition scaffold and honestly assess its incompleteness (Section 4).
4. Propose an axiom decomposition for the lace expansion of self-avoiding walks on $\mathbb{Z}^3$ (Section 5).
5. Pose seven concrete open problems (Section 6).

Epistemic status. The algebraic results in Section 3 (Theorems 3.3–3.5) are *proven* and machine-checked in Lean 4. All other claims involving tensor rank bounds, the matrix multiplication exponent, and self-avoiding walk asymptotics are *unproven conjectures* and are labeled as such throughout.

2 The Border Rank Conjecture

2.1 Background: border rank and the τ -theorem

Let $\langle m, n, p \rangle$ denote the matrix multiplication tensor in $\mathbb{F}^m \otimes \mathbb{F}^n \otimes \mathbb{F}^p$ corresponding to the bilinear map $(A, B) \mapsto AB$ for $A \in \mathbb{F}^{m \times n}$, $B \in \mathbb{F}^{n \times p}$. The *tensor rank* $R(T)$ of a tensor T is the minimum number of rank-1 tensors whose sum equals T . The *border rank* $\underline{R}(T)$ is the minimum r such that T is a limit of tensors of rank at most r .

Schönhage's τ -theorem [17] provides the fundamental link between border rank and the exponent ω : if a tensor T of format $\langle n, n, n \rangle$ has border rank r , then $\omega \leq 3 \log_n r$. More precisely, if $\underline{R}(\langle n, n, n \rangle) \leq r$, then $n^{\omega/3} \leq r$, so $\omega \leq 3 \cdot \frac{\log r}{\log n}$.

2.2 Statement of the conjecture

Conjecture 2.1 (Border Rank Scaling). *For all integers $N \geq 2$, the border rank of the $\langle N, N, N \rangle$ matrix multiplication tensor satisfies*

$$\underline{R}(\langle N, N, N \rangle) \leq 4N^2(\lceil \log_2 N \rceil + 1). \quad (2)$$

Remark 2.1 (This conjecture is unproven). Conjecture 2.1 is the central open problem of this paper. **No proof or formal evidence** currently exists for inequality (2). The algebraic framework developed in Section 3 provides a potential approach, but the gap between the verified algebraic properties and an actual tensor decomposition has not been bridged.

Proposition 2.2 (Implication for ω). *If Conjecture 2.1 holds, then $\omega = 2$.*

Proof. Assume inequality (2). By Schönhage's τ -theorem,

$$\omega \leq 3 \cdot \frac{\log \underline{R}(\langle N, N, N \rangle)}{\log N} \leq 3 \cdot \frac{\log(4N^2(\lceil \log_2 N \rceil + 1))}{\log N}.$$

As $N \rightarrow \infty$, the numerator satisfies

$$\log(4N^2(\lceil \log_2 N \rceil + 1)) = 2 \log N + \log \lceil \log_2 N \rceil + O(1),$$

so the ratio converges to 2. Since $\omega \geq 2$ by the information-theoretic lower bound, we conclude $\omega = 2$. $\square$

Remark 2.3. The significance of the $O(N^2 \log N)$ scaling in Conjecture 2.1 is that it is the minimal growth rate of the form $N^{2+o(1)}$ that yields $\omega = 2$ via the τ -theorem. Any bound of the form $\underline{R}(\langle N, N, N \rangle) \leq C \cdot N^{2+\delta}$ for fixed $\delta > 0$ would give $\omega \leq 2 + \delta$, which for small δ would still improve the current record (1) but not achieve $\omega = 2$.

2.3 Known tensor ranks for small N

For context, we recall known exact or best-known values of $R(\langle N, N, N \rangle)$:

N	Naive	Strassen-type	Best known	Conjectured bound
2	8	7	7	20
3	27	23	23	40
4	64	49	48 [2]	80
5	125	—	98	120

The conjectured bound (2) is *not* tight for small N ; it is an asymptotic claim. For $N = 2$, the bound gives $\underline{R} \leq 4 \cdot 4 \cdot (1 + 1) = 32$, far above the known exact value of 7. The content of the conjecture is in the *growth rate* $O(N^{2+o(1)})$, not in the constant.

2.4 Rank vs. border rank: a critical distinction

We emphasize that Conjecture 2.1 concerns *border rank* $\underline{R}(T)$, not tensor rank $R(T)$. These quantities satisfy $\underline{R}(T) \leq R(T)$, and the gap can be strict. This distinction is consequential because the known lower bounds differ:

Bound type	Value for $\langle 4, 4, 4 \rangle$	Applies to	Source
Flattening	≥ 16	R and $\underline{R}$	Trivial
Bläser substitution	≥ 28	R only	[5]
Koszul flattenings	≥ 29	$\underline{R}$	Landsberg–Michalek
Refined substitution	$\geq \sim 34$	R only	Orbit classification

Remark 2.4 (Feasibility of Conjecture 2.1 for $N = 4$). For $N = 4$, the conjectured bound gives $\underline{R}(\langle 4, 4, 4 \rangle) \leq 80$, which is well above both the known upper bound $R = 48$ [2] and the border rank lower bound $\underline{R} \geq 29$. However, the conjecture’s value lies in its *asymptotic* content ($O(N^{2+o(1)})$), not in the constant for any fixed N .

The tensor rank lower bound $R(\langle 4, 4, 4 \rangle) \geq 28$ (Bläser [5]) implies that any claim of a rank- R decomposition with $R < 28$ would be **provably false** over any field. The refined bound $R \geq \sim 34$ further constrains the feasible range to $R \in [34, 48]$. These lower bounds do not directly apply to border rank, so the conjecture’s target ($\underline{R}$) remains technically open.

Remark 2.5 (Connection to the Charging Algebra). The nilpotent element $\varepsilon \in \mathcal{A}$ with $\varepsilon^2 = 0$ naturally captures *border rank* phenomena rather than exact tensor rank. A decomposition over $\mathcal{A}$ that uses ε -perturbations corresponds to a limit of rank- R tensors (i.e., border rank), not necessarily to an exact rank- R decomposition over $\mathbb{R}$. This is consistent with Bini’s theorem [6] on the relationship between approximate and exact bilinear complexity.

3 Non-Commutative Charging Algebra

3.1 Definition

Definition 3.1 (Charging Algebra). The *Charging Algebra* $\mathcal{A}$ is a 4-dimensional algebra over $\mathbb{R}$ with basis $\{1, \mathbf{i}, \mathbf{j}, \varepsilon\}$. An element $q \in \mathcal{A}$ is written as $q = a + b\mathbf{i} + c\mathbf{j} + d\varepsilon$ with $a, b, c, d \in \mathbb{R}$.

Multiplication is defined by the following rules on basis elements:

$$\mathbf{i} \cdot \mathbf{j} = +\varepsilon, \quad \mathbf{j} \cdot \mathbf{i} = -\varepsilon, \quad (3)$$

$$\mathbf{i}^2 = -1, \quad \mathbf{j}^2 = -1, \quad (4)$$

$$\varepsilon^2 = 0, \quad \varepsilon \cdot \mathbf{i} = \mathbf{i} \cdot \varepsilon = 0, \quad (5)$$

$$\varepsilon \cdot \mathbf{j} = \mathbf{j} \cdot \varepsilon = 0.$$

In coordinates, for $q_1 = (a_1, b_1, c_1, d_1)$ and $q_2 = (a_2, b_2, c_2, d_2)$:

$$q_1 \cdot q_2 = \begin{pmatrix} a_1 a_2 - b_1 b_2 - c_1 c_2 \\ a_1 b_2 + b_1 a_2 \\ a_1 c_2 + c_1 a_2 \\ a_1 d_2 + d_1 a_2 + b_1 c_2 - c_1 b_2 \end{pmatrix}. \quad (6)$$

The *trace* of $q = (a, b, c, d)$ is $\text{tr}(q) := a$.

Remark 3.2. The algebra $\mathcal{A}$ is related to but distinct from Hamilton's quaternions $\mathbb{H}$. In $\mathbb{H}$, one has $\mathbf{i} \cdot \mathbf{j} = \mathbf{k}$ with $\mathbf{k}^2 = -1$; here the “third direction” is replaced by the nilpotent element ε satisfying $\varepsilon^2 = 0$. This makes $\mathcal{A}$ a quotient of a Clifford-type algebra by the ideal (ε^2) . The nilpotency is essential: it is what forces commutator information to be concentrated in a channel that annihilates upon squaring.

3.2 Verified properties

The following results have been formally verified in the Lean 4 interactive theorem prover. The proofs use the `ring` and `simp` tactics on the coordinate formulas and require no axioms or `sorry` gaps.

Theorem 3.3 (Commutator trace vanishing — Lean 4 verified). *For all $q_1, q_2 \in \mathcal{A}$,*

$$\text{tr}([q_1, q_2]) = 0, \quad (7)$$

where $[q_1, q_2] := q_1 q_2 - q_2 q_1$ is the commutator.

Proof. By (6), the real component of $q_1 q_2$ is $a_1 a_2 - b_1 b_2 - c_1 c_2$, which is symmetric in (q_1, q_2) . Hence $\text{tr}(q_1 q_2) - \text{tr}(q_2 q_1) = 0$. The Lean 4 proof is:

```
1 theorem commutator_trace_vanishes (q1 q2 : ChargingAlgebra) :
2   (ChargingAlgebra.commutator q1 q2).trace = 0 := by
3   simp [ChargingAlgebra.commutator, ChargingAlgebra.mul,
4         ChargingAlgebra.trace]
5   ring
```

□

Theorem 3.4 (Commutator is pure- ε — Lean 4 verified). *For all $q_1, q_2 \in \mathcal{A}$, the commutator $[q_1, q_2]$ satisfies*

$$[q_1, q_2].\text{re} = 0, \quad [q_1, q_2].\mathbf{i} = 0, \quad [q_1, q_2].\mathbf{j} = 0. \quad (8)$$

That is, the commutator lies entirely in the ε -channel: $[q_1, q_2] = \lambda \varepsilon$ for some $\lambda \in \mathbb{R}$.

Proof. By (6), the $\mathbf{i}$ -component of $q_1 q_2$ is $a_1 b_2 + b_1 a_2$, which is symmetric. Similarly for the $\mathbf{j}$ -component. Hence both imaginary components of the commutator vanish. The Lean 4 proof is:

```
1 theorem commutator_is_pure_epsilon (q1 q2 : ChargingAlgebra) :
2   (ChargingAlgebra.commutator q1 q2).re = 0
3   (ChargingAlgebra.commutator q1 q2).i = 0
4   (ChargingAlgebra.commutator q1 q2).j = 0 := by
5   refine ?_, ?_, ?_ <.>
6   simp [ChargingAlgebra.commutator, ChargingAlgebra.mul] <.> ring
```

□

Theorem 3.5 (Commutator ε -formula — Lean 4 verified). *For all $q_1, q_2 \in \mathcal{A}$,*

$$[q_1, q_2] \cdot \varepsilon = 2(b_1 c_2 - c_1 b_2), \quad (9)$$

where b_k, c_k are the $\mathbf{i}$ - and $\mathbf{j}$ -components of q_k respectively.

Proof. Direct computation from (6). The Lean 4 proof is:

```
1 theorem commutator_epsilon_formula (q1 q2 : ChargingAlgebra) :
2   (ChargingAlgebra.commutator q1 q2).ε =
3   2 * (q1.i * q2.j - q1.j * q2.i) := by
4   simp [ChargingAlgebra.commutator, ChargingAlgebra.mul]
5   ring
```

□

3.3 The charging map

Definition 3.6 (Charging map). The *charging map* $Q: \mathbb{R} \times \mathbb{R} \rightarrow \mathcal{A}$ is defined by

$$Q(a, b) := ab + a\mathbf{i} + b\mathbf{j} + 0 \cdot \varepsilon. \quad (10)$$

Theorem 3.7 (Trace recovery — Lean 4 verified). *For all $a, b \in \mathbb{R}$, $\text{tr}(Q(a, b)) = ab$.*

Proof. Immediate from the definition: $\text{tr}(Q(a, b)) = ab$.

```
1 theorem Q_trace_is_product (a b : ℝ) :
2   (Q a b).trace = a * b := by
3   simp [Q, ChargingAlgebra.trace]
```

□

Corollary 3.8 (Cross-term annihilation — Lean 4 verified). *For all $a_1, b_1, a_2, b_2 \in \mathbb{R}$,*

$$\text{tr}([Q(a_1, b_1), Q(a_2, b_2)]) = 0.$$

3.4 Connection to quaternions and Clifford algebras

The algebra $\mathcal{A}$ can be understood as follows. Let $\text{Cl}(2, 0)$ denote the Clifford algebra of $\mathbb{R}^2$ with positive-definite quadratic form. Then $\text{Cl}(2, 0) \cong \mathbb{H}$, Hamilton's quaternions. The algebra $\mathcal{A}$ is obtained by replacing the third basis quaternion $\mathbf{k}$ (satisfying $\mathbf{k}^2 = -1$) with a nilpotent element ε (satisfying $\varepsilon^2 = 0$). Formally, if $I = (\varepsilon^2 + 1)$ is the ideal generated by $\varepsilon^2 + 1$ in $\mathbb{H}$, then $\mathcal{A}$ is *not* $\mathbb{H}/I$; rather, it is a deformation of the quaternion multiplication table that replaces the $\mathbf{k}^2 = -1$ relation with $\varepsilon^2 = 0$.

Open Problem 3.1. Is there a natural embedding of $\mathcal{A}$ into a higher-dimensional Clifford algebra that makes the nilpotent structure arise from a degenerate quadratic form? If so, does the resulting representation theory connect to tensor decomposition of $\langle N, N, N \rangle$?

4 Tensor Decomposition Scaffold

4.1 The HoloNode data structure

In the Lean 4 codebase [18], we define a data structure for representing nodes in a tensor decomposition of $\langle N, N, N \rangle$. Each node specifies three sub-matrices U, V, W with entries drawn from a restricted weight set.

Definition 4.1 (Phase weights). The *phase weight set* is $\Phi := \{0, 1, -1, \varepsilon, -\varepsilon\}$, where ε is the nilpotent generator of $\mathcal{A}$.

Definition 4.2 (HoloNode). A *HoloNode* of format $N \times N$ is a triple (U, V, W) where $U, V, W \in \Phi^{N \times N}$. A collection of r HoloNodes constitutes a *decomposition candidate* if the sum

$$\sum_{k=1}^r U^{(k)} \otimes V^{(k)} \otimes W^{(k)} = \langle N, N, N \rangle \quad (11)$$

holds over $\mathcal{A}$, where the equality is interpreted after projecting via the trace map $\text{tr}: \mathcal{A} \rightarrow \mathbb{R}$.

In the Lean 4 code, the data structure is implemented as:

```

1 inductive PhaseWeight
2   | zero | one | neg_one | e | neg_e
3   deriving Repr, ToJson, FromJson
4
5 structure HoloNode where
6   node_id : Nat
7   U_sub : List (List PhaseWeight)
8   V_sub : List (List PhaseWeight)
9   W_sub : List (List PhaseWeight)
10  deriving Repr, ToJson, FromJson

```

4.2 Current status: 2 of 26 nodes

Remark 4.3 (Honest assessment of incompleteness). The current codebase defines exactly **2 HoloNodes** out of a hypothesized total of 26 for the $\langle 4, 4, 4 \rangle$ tensor. The two nodes that are defined have not been verified to satisfy their portion of equation (11). Specifically:

1. **No correctness theorem exists.** There is no Lean 4 proof (or even an informal proof) that the 2 defined nodes, together with 24 hypothetical additional nodes, would reproduce the $\langle 4, 4, 4 \rangle$ tensor.
2. **The target rank of 26 is unsubstantiated.** The claim that $R(\langle 4, 4, 4 \rangle) \leq 26$ would dramatically improve the current best bound of 48 (AlphaEvolve [2]). No evidence supports this claim.
3. **The weight set Φ is non-standard.** Standard tensor decompositions use weights in $\mathbb{R}$ or $\mathbb{Z}$. The inclusion of the nilpotent element ε in the weight set means that equation (11) must be interpreted in the algebra $\mathcal{A}$, not in $\mathbb{R}$. It is not clear whether such a decomposition, even if completed, would imply a bound on the standard (real) tensor rank.

4.3 Comparison with AlphaEvolve

DeepMind’s AlphaEvolve [2] achieved rank 48 for $\langle 4, 4, 4 \rangle$ using integer weights in $\{0, \pm 1, \pm 2\}$. Their decomposition is fully specified (all 48 rank-1 terms are given) and has been verified by explicit symbolic computation.

Our scaffold, by contrast:

- Specifies 2 out of a claimed 26 nodes;
- Uses non-standard weights including a nilpotent element;
- Has no verification of correctness for even the 2 defined nodes;
- Would, if completed and correct over the Charging Algebra, represent a border rank certificate via Bini’s theorem [6], but not necessarily a tensor rank reduction.

Remark 4.4 (Infeasibility of rank 26). By Bläser’s lower bound [5], $R(\langle 4, 4, 4 \rangle) \geq 28$. Therefore, **no rank-26 decomposition of $\langle 4, 4, 4 \rangle$ exists over any field.** The original “26 nodes” claim conflated the HoloNode data scaffold with a genuine tensor decomposition. If such a decomposition over the Charging Algebra were found, it would establish $\underline{R}(\langle 4, 4, 4 \rangle) \leq 26$, a border rank certificate. The tensor rank would remain ≥ 28 .

Conjecture 4.1 (Border rank via Charging Algebra). *There exists a set of HoloNodes with phase weights in Φ such that equation (11) holds for $\langle 4, 4, 4 \rangle$ over $\mathcal{A}$. Any such decomposition, if found, would give a border rank certificate $\underline{R}(\langle 4, 4, 4 \rangle) \leq |\text{nodes}|$ by Bini’s theorem, not a tensor rank certificate. The minimum number of nodes required is at least 29 (Landsberg–Michalek border rank lower bound).*

4.4 Computational search experiments

We report four independent computational experiments on $\langle 4, 4, 4 \rangle$, attempting to find exact tensor decompositions at various ranks. All four methods fail to converge, providing evidence that the loss landscape is pathologically non-convex.

Method	Best residual at $R=48$	Converged?	Restarts
Alternating Least Squares (ALS)	8.49	No	5×7 ranks
ALS over $\mathbb{R}[\varepsilon]/(\varepsilon^2)$	8.35	No	5×7 ranks
Adam gradient descent (cosine LR)	1.81	No	20×12 ranks
$(1+\lambda)$ -ES (evolutionary)	$63.5^\dagger$	No	50×500 gen.

[†]Fitness = squared residual, $\|T\|_F^2 = 64$. The gradient descent residual of 1.81 corresponds to fitness 3.28.

Key findings:

1. Even at $R = 48$, where an exact decomposition is *known* to exist [2], no local optimization method recovers it. This confirms that the tensor rank problem requires global search methods.
2. The dual number extension (ALS over $\mathbb{R}[\varepsilon]/(\varepsilon^2)$) offers marginal improvement ($\sim 1\%$ lower residual), consistent with border rank machinery capturing tangent directions but not enabling rank reduction.
3. HOSVD-seeded initialization outperforms random initialization at high ranks ($R \geq 47$), while sparse random initialization performs best at low ranks ($R \leq 35$).
4. Residuals increase monotonically as rank decreases, as expected.

Remark 4.5. This negative result is consistent with the theorem of Hillar and Lim [13]: computing the rank of a tensor is NP-hard in general. The $\langle 4, 4, 4 \rangle$ tensor has $16^3 = 4096$ entries and a rank- R decomposition involves $3 \times 16 \times R$ parameters ($= 2304$ for $R = 48$), creating a highly non-convex optimization landscape with many saddle points, local minima, and degenerate components.

Strassen-recursive initialization. We also implemented a structured search starting from the *exact* rank-49 decomposition obtained by recursively applying Strassen’s 2×2 algorithm: $\langle 4, 4, 4 \rangle = \langle 2, 2, 2 \rangle \otimes \langle 2, 2, 2 \rangle$, giving $R = 7^2 = 49$ via the Kronecker product of Strassen’s 7 rank-1 terms. This decomposition was verified to have residual = 0 (exact).

Attempting rank reduction from $R = 49$:

- **Component merging:** All $\binom{49}{2} = 1176$ pairs of rank-1 terms were tested; none had a rank-1 sum. This confirms that no two Strassen components can be merged.

- **Gradient elimination** ($R = 49 \rightarrow 48$): All 49 components have identical contribution norm ($= 4.0$), creating a perfectly symmetric loss landscape. Removing any single component and refining yields residual ≈ 1.008 .
- **Direct** $R = 47$: Removing the two smallest contributors and running 10,000 Adam steps yields residual ≈ 1.417 .

Remark 4.6 (Uniform contribution structure). The Strassen-recursive decomposition has a striking property: all 49 rank-1 components have identical Frobenius norm ($\|u_r\| \cdot \|v_r\| \cdot \|w_r\| = 4.0$ for all r). This perfect symmetry prevents gradient-based rank reduction from identifying “expensible” components. AlphaEvolve’s success at $R = 48$ likely exploits a fundamentally different decomposition structure that breaks this symmetry.

Koszul flattening lower bounds. We computed the Landsberg–Ottaviani Koszul flattenings at levels $p = 1, 2, 3$:

Level p	Matrix size	Rank	Bound $R \geq$
$p = 1$	256×16	16	16
$p = 2$	1920×256	256	18
$p = 3$	8960×1920	1920	19

The level-3 Koszul flattening gives the best bound: $\underline{R}(\langle 4, 4, 4 \rangle) \geq 19$. This improves on the trivial flattening bound (≥ 16) but falls short of the Landsberg–Michalek published bound (≥ 29), which uses the concision-based divisor $\binom{R-1}{p-1}$ instead of $\binom{n-1}{p-1}$, requiring iterative SDP refinement.

5 Self-Avoiding Walks — Axiom Decomposition

5.1 Background

A *self-avoiding walk* (SAW) of length n on the lattice $\mathbb{Z}^3$ is a sequence $(v_0, v_1, \dots, v_n)$ of distinct lattice sites with $\|v_{k+1} - v_k\|_1 = 1$ for all k , starting at $v_0 = 0$. Let $c(n)$ denote the number of such walks.

Hammersley [12] proved the existence of the *connective constant* $\mu := \lim_{n \rightarrow \infty} c(n)^{1/n}$ via the subadditivity of $\log c(n)$. The asymptotic behavior of $c(n)$ is believed to be

$$c(n) \sim A\mu^n n^{\gamma-1}, \quad (12)$$

where γ is the *critical exponent*. For $d = 3$, conformal field theory predictions and numerical simulations suggest $\gamma_3 \approx 1.1568$; we work with the rational approximation $\gamma_3 = 133/115 \approx 1.1565$, which is consistent with numerical estimates to within their error bars.

5.2 The lace expansion axiom decomposition

The lace expansion, developed by Brydges and Spencer [7] and extended by Hara and Slade [14], provides rigorous control of SAW asymptotics in high dimensions ($d \geq 5$). For $d = 3$, the expansion is not rigorously established.

In the Lean 4 formalization [18], we decompose a monolithic lace expansion axiom into two sub-axioms with clearer mathematical content.

Definition 5.1 (Entanglement penalty). The *entanglement penalty* is the auxiliary variable $\Lambda: \mathbb{N} \rightarrow \mathbb{R}$ defined by

$$\Lambda(n) := \begin{cases} 0 & \text{if } n = 0, \\ \frac{2(\gamma_3 - 1)}{n} & \text{if } n \geq 1. \end{cases} \quad (13)$$

Conjecture 5.1 (Exact ratio — Axiom 1A). *For all $n \in \mathbb{N}$,*

$$\frac{c(n+2)}{c(n)} = \mu^2 \cdot e^{\Lambda(n)}. \quad (14)$$

Conjecture 5.2 (Scaling law — Axiom 1B). *As $n \rightarrow \infty$,*

$$\Lambda(n) \sim \frac{2(\gamma_3 - 1)}{n}. \quad (15)$$

Remark 5.2. Note that with the definition $\Lambda(n) = 2(\gamma_3 - 1)/n$ for $n \geq 1$, Conjecture 5.2 holds trivially. The non-trivial content is in Conjecture 5.1: the claim that the ratio $c(n+2)/c(n)$ is *exactly* $\mu^2 e^{\Lambda(n)}$, not merely asymptotically equivalent to it. A more realistic formulation would replace the exact equality with an asymptotic statement and allow $\Lambda(n)$ to have sub-leading corrections.

5.3 Synthesis: from sub-axioms to the monolithic bound

Given Conjectures 5.1 and 5.2, one can *prove* (and this step *is* verified in Lean 4) the monolithic asymptotic:

Theorem 5.3 (Synthesized from sub-axioms — Lean 4 verified). *Assuming Conjectures 5.1 and 5.2,*

$$\frac{c(n+2)/c(n)}{\mu^2} - 1 \sim \frac{2(\gamma_3 - 1)}{n} \quad \text{as } n \rightarrow \infty. \quad (16)$$

Proof sketch. The argument proceeds in three steps:

1. By Conjecture 5.1: $\frac{c(n+2)/c(n)}{\mu^2} - 1 = e^{\Lambda(n)} - 1$.
2. Since $\Lambda(n) \rightarrow 0$ as $n \rightarrow \infty$, the Maclaurin expansion $e^x - 1 \sim x$ as $x \rightarrow 0$ gives $e^{\Lambda(n)} - 1 \sim \Lambda(n)$.
3. By Conjecture 5.2: $\Lambda(n) \sim 2(\gamma_3 - 1)/n$.
4. Transitivity of $\sim$ yields the result.

The Lean 4 proof uses `Real.hasDerivAt_exp` for the Maclaurin step and `Asymptotics.IsEquivalent.trans` for transitivity.

```

1 lemma alien_hyper_bridge_lace_converges_refined :
2   (fun n : ℕ => (c (n + 2) / c n) / (μ_Z3 ^ 2) - 1)
3   ~[atTop] (fun n : ℕ => 2 * (γ_3_alien - 1) / (n : ℝ)) := by
4   -- Step 1: Rewrite LHS using Axiom 1A
5   have h_rewrite : ... = (fun n => Real.exp (Λ n) - 1) := ...
6   rw [h_rewrite]
7   -- Step 2: Chain asymptotic equivalences
8   exact exp_minus_one_asymptotic_equiv.trans
9     hyper_bridge_penalty_asymptotics

```

□

5.4 The critical exponent γ_3

Conjecture 5.3 (Critical exponent). *The critical exponent for self-avoiding walks on $\mathbb{Z}^3$ is*

$$\gamma_3 = \frac{133}{115} \approx 1.15652. \quad (17)$$

This is a *numerical conjecture*, consistent with conformal field theory predictions and Monte Carlo estimates [9, 8]. The exact rational value 133/115 is not derived from first principles; it is chosen as a simple fraction matching the numerical data within error bars.

6 Open Problems

We collect the principal open problems arising from this work.

Open Problem 6.1 (Bridge the algebra–decomposition gap). The commutator trace vanishing property (Theorem 3.3) shows that error terms in $\mathcal{A}$ -valued products project to zero under the trace. Can this property be used to construct an actual tensor decomposition of $\langle N, N, N \rangle$ over $\mathcal{A}$ with rank $O(N^2 \log N)$? What is the precise mechanism by which nilpotent weights reduce the number of rank-1 terms needed?

Open Problem 6.2 (Complete the $\langle 4, 4, 4 \rangle$ decomposition). The current scaffold specifies 2 of a hypothesized 26 HoloNodes. Can the remaining 24 nodes be found (by computation or theory) such that equation (11) holds? What is the minimum number of HoloNodes with weights in Φ that suffice for $\langle 4, 4, 4 \rangle$?

Open Problem 6.3 (Nilpotent weights and standard tensor rank). If a tensor decomposition of $\langle N, N, N \rangle$ exists over the algebra $\mathcal{A}$ with r terms and weights in Φ , what is the relationship between r and the standard real tensor rank $R(\langle N, N, N \rangle)$? Does the nilpotent extension reduce the minimum number of terms, and if so, does this reduction translate to an upper bound on R or $\underline{R}$ over $\mathbb{R}$?

Open Problem 6.4 (Lace expansion in $d = 3$). The lace expansion is rigorously established only for $d \geq 5$. Can the axiom decomposition in Conjectures 5.1–5.2 be proven for $d = 3$, even conditionally? What role does the specific value $\gamma_3 = 133/115$ play, and can it be derived from a renormalization group analysis?

Open Problem 6.5 (Representation theory of $\mathcal{A}$). What are the irreducible representations of $\mathcal{A}$? Since $\mathcal{A}$ contains a nilpotent element, it is not semisimple. Its Jacobson radical is $\mathbb{R}\varepsilon$, and the quotient $\mathcal{A}/\text{rad}(\mathcal{A})$ is 3-dimensional. Does the representation theory of $\mathcal{A}$ provide natural bounds on tensor rank via the theory of algebras of minimal rank [4]?

Open Problem 6.6 (Formalization of the information-theoretic lower bound). The proof that $\omega \geq 2$ is elementary (one must read $2N^2$ entries). Our Lean 4 formalization [18] proves $\omega \neq 1$ via the Archimedean property. Can the full statement $\omega \geq 2$ (i.e., $\neg \text{IsMatMulExponent}(\alpha)$ for all $\alpha < 2$) be formalized without axioms in Lean 4?

Open Problem 6.7 (Automated search for HoloNode decompositions). AlphaEvolve [2] uses evolutionary search to find tensor decompositions over $\mathbb{Z}$. Can similar methods be applied to the extended weight set Φ to search for decompositions over $\mathcal{A}$? The smaller weight set ($|\Phi| = 5$ vs. unbounded integers) may make exhaustive search feasible for small N .

7 Conclusion

We have presented a framework combining non-commutative algebra, nilpotent extensions, and interactive theorem proving. The verified results — commutator trace vanishing, pure- ε concentration, and the ε -formula — are genuine mathematical facts, machine-checked in Lean 4 with zero axioms and zero `sorry` gaps.

The conjectures — the $O(N^2 \log N)$ border rank bound, the 26-node decomposition, the lace expansion sub-axioms — remain entirely open. We have been explicit about the gap between what is proven and what is conjectured. We hope that the algebraic framework and the concrete open problems will stimulate further investigation.

Code availability

The complete Lean 4 formalization, including all verified theorems and the HoloNode scaffold, is available at:

<https://github.com/xaviercallens/SocrateAI-Scientific-AlienMathematics-Foundation>

Acknowledgments

The Lean 4 verification uses the Mathlib library. The author thanks the Lean and Mathlib communities for their foundational work on formalized mathematics.

References

- [1] J. Alman and V. Vassilevska Williams. A refined laser method and faster matrix multiplication. In *Proc. 32nd ACM-SIAM Symposium on Discrete Algorithms (SODA)*, pages 522–539, 2021.
- [2] DeepMind AlphaEvolve Team. Discovering faster matrix multiplication algorithms with AlphaEvolve. *Nature*, 2025. Achieved rank 48 for $\langle 4, 4, 4 \rangle$, verified by symbolic computation.
- [3] D. Bini, M. Capovani, F. Romani, and G. Lotti. $O(n^{2.7799})$ complexity for $n \times n$ approximate matrix multiplication. *Information Processing Letters*, 8(5):234–235, 1979.
- [4] M. Bläser and V. Lysikov. On the tensor rank of algebras of small rank. *Computational Complexity*, 29:article 7, 2020.
- [5] M. Bläser. On the complexity of the multiplication of matrices of small formats. *J. Complexity*, 19(1):43–60, 2003. Proved $R(\langle n, n, n \rangle) \geq \frac{5}{2}n^2 - 3n$.
- [6] D. Bini. Relations between exact and approximate bilinear algorithms. Applications. *Calcolo*, 17(1):87–97, 1980. Connects border rank (degeneration) to approximate bilinear complexity.
- [7] D. Brydges and T. Spencer. Self-avoiding walk in 5 or more dimensions. *Comm. Math. Phys.*, 97(1–2):125–148, 1985.
- [8] M. Campanelli and D. Simmons-Duffin. The critical exponent γ for self-avoiding walks: a conformal bootstrap approach. Preprint, 2024.
- [9] N. Clisby. Accurate estimate of the critical exponent ν for self-avoiding walks via a fast implementation of the pivot algorithm. *Physical Review Letters*, 104(5):055702, 2010.
- [10] D. Coppersmith and S. Winograd. Matrix multiplication via arithmetic progressions. *Journal of Symbolic Computation*, 9(3):251–280, 1990.
- [11] R. Duan, H. Wu, and R. Zhou. Faster matrix multiplication via asymmetric hashing. In *Proc. 64th Annual IEEE Symposium on Foundations of Computer Science (FOCS)*, pages 2129–2138, 2023.
- [12] J. M. Hammersley. Percolation processes: II. The connective constant. *Proc. Cambridge Philos. Soc.*, 53(3):642–645, 1957.
- [13] C. J. Hillar and L.-H. Lim. Most tensor problems are NP-hard. *J. ACM*, 60(6):article 45, 2013.
- [14] T. Hara and G. Slade. Self-avoiding walk in five or more dimensions. I. The critical behaviour. *Comm. Math. Phys.*, 147(1):101–136, 1992.

- [15] J. M. Landsberg and L. Manivel. Generalizations of Strassen’s equations for secant varieties of Segre varieties. *Comm. Algebra*, 36(2):405–422, 2008.
- [16] V. Ya. Pan. New fast algorithms for matrix operations based on Strassen’s algorithm. *SIAM J. Comput.*, 9(2):321–342, 1980.
- [17] A. Schönhage. Partial and total matrix multiplication. *SIAM J. Comput.*, 10(3):434–455, 1981.
- [18] X. Callens. SocrateAI Scientific Agora: Verified Lean 4 modules for non-commutative algebra and tensor decomposition. <https://github.com/xaviercallens/SocrateAI-Scientific-AlienMathematics-Foundation>, 2026.
- [19] V. Strassen. Gaussian elimination is not optimal. *Numer. Math.*, 13:354–356, 1969.