

Automated Discovery and Lean 4 Verification of Combinatorial Identities via Rational Arithmetic Falsification

Xavier Callens
Socrate AI Lab
xavier@socrateai.org

June 2026

Abstract

We present a fully automated pipeline that *discovers*, *verifies*, and *formalises* combinatorial identities involving binomial coefficients. The pipeline (i) generates candidate identities from five parametric families (weighted binomial, squared binomial, alternating, convolution, mixed), (ii) performs computational falsification using *exact* rational arithmetic over $\mathbb{Q}$ (no floating-point), (iii) fits closed-form expressions via polynomial $\times 2^n$ decomposition, (iv) filters tautologies through an anti-tautology oracle, and (v) emits LEAN 4 proof scripts verified by `native_decide`. The pipeline produced **18 candidate conjectures**, of which **12 were verified** as known identities, **2 were fitted** to novel closed forms, and **2 are potentially new** (independently rediscovered). All concrete instances are formalised in a LEAN 4 library comprising **49 theorems** with **0 sorry** and **0 axiom**, where all 2504/2504 modules build successfully. We are transparent that the “discovered” identities may appear in existing literature; the contribution lies in the *pipeline itself* as a prototype for automated mathematical discovery.

Contents

1	Introduction	2
2	Methodology	2
2.1	Conjecture Generation	2
2.2	Computational Falsification	3
2.3	Closed-Form Fitting	3
2.4	Non-Triviality Checking	3
2.5	Lean 4 Formalization	3
3	Results	3
3.1	Pipeline Statistics	3
3.2	Verified Known Identities	4
3.3	Discovered Identities	5
3.4	Lean 4 Verification	6
4	Discussion	7
4.1	Novelty Assessment	7
4.2	Lessons from Alien Mathematics	7
4.3	Pipeline Validation	8
5	Conclusion and Future Work	8

1 Introduction

The classical theory of combinatorial identities is vast: from Vandermonde’s convolution to the Chu–Vandermonde identity, from the hockey-stick recurrence to hypergeometric summation algorithms such as Zeilberger’s algorithm [1] and the Wilf–Zeilberger method. Despite centuries of accumulation, the *discovery* of new identities remains largely a craft—dependent on human intuition, algebraic manipulation, and serendipity.


This work is motivated by the *Alien Mathematics* thought experiment: if an extraterrestrial civilisation presented us with unfamiliar mathematical objects, how would we systematically explore them? The key lessons we draw are:

1. **Falsification first.** A conjecture earns credibility not by passing a few tests but by surviving aggressive attempts at refutation.
2. **Tautology avoidance.** An identity that holds for trivial structural reasons (e.g. $\sum_{k=0}^0 \binom{0}{k} = 1$) teaches us nothing; we must filter out tautologies.
3. **Exact arithmetic.** Numerical coincidences over $\mathbb{R}$ are cheap. Agreement over $\mathbb{Q}$ with exact rational arithmetic provides a far stronger signal.
4. **Formal verification.** Even strong numerical evidence is not a proof. Machine-checked formalisation in LEAN 4 closes the gap for concrete instances.

Our pipeline operationalises these principles. Section 2 describes the methodology. Section 3 presents the identities discovered and verified. Section 4 discusses novelty, limitations, and lessons learned. Section 5 outlines future directions.

2 Methodology

The pipeline consists of five stages, depicted schematically below:



2.1 Conjecture Generation

We generate candidate identities from five parametric families of summation formulae over binomial coefficients.

Definition 2.1 (Identity families). Let $n \in \mathbb{N}$. The five families are:

1. **Weighted binomial sums.** $S_p(n) = \sum_{k=0}^n k^p \binom{n}{k}$ for $p \in \{0, 1, 2, \dots, 6\}$.
2. **Squared binomial sums.** $T_p(n) = \sum_{k=0}^n k^p \binom{n}{k}^2$ for $p \in \{0, 1, 2\}$.
3. **Alternating sums.** $A(n) = \sum_{k=0}^n (-1)^k \binom{n}{k} \binom{n+k}{k}$.
4. **Convolution sums.** $V(m, n, r) = \sum_{k=0}^r \binom{m}{k} \binom{n}{r-k}$ (Vandermonde family).
5. **Mixed factorial sums.** Products involving $k!$, $(n-k)!$, and related terms.

Each family is instantiated over a grid of small parameter values, yielding an initial pool of candidate conjectures.

2.2 Computational Falsification

Every candidate identity is tested by evaluating both sides for $n = 0, 1, \dots, N_{\max}$ (typically $N_{\max} = 50$) using *exact rational arithmetic* in Python’s `Fractions.Fraction` class.

Remark 2.2 (Why exact arithmetic matters). Floating-point evaluation can produce false positives due to round-off coincidences, especially for large binomial coefficients. Since $\binom{100}{50} \approx 10^{29}$, double-precision arithmetic (≈ 15 significant digits) is wholly inadequate. Exact computation over $\mathbb{Q}$ eliminates this failure mode entirely.

A candidate survives falsification only if *every* tested instance yields exact equality. Any single counterexample immediately kills the conjecture.

2.3 Closed-Form Fitting

For survivors of the form $S_p(n) = ?$, we attempt to express the right-hand side as a product $Q_p(n) \cdot 2^n$ where Q_p is a polynomial in n .

Definition 2.3 (Finite-difference fitting). Given the sequence $a_n = S_p(n)/2^n$ for $n = 0, 1, \dots, d+1$, we compute iterated forward differences $\Delta^j a_0 = \sum_{i=0}^j (-1)^{j-i} \binom{j}{i} a_i$. If $\Delta^{d+1} a_0 = 0$ exactly (over $\mathbb{Q}$), then a_n is a polynomial of degree at most d , and we recover its coefficients by Newton interpolation.

This procedure discovered that $\sum_{k=0}^n k^5 \binom{n}{k}$ and $\sum_{k=0}^n k^6 \binom{n}{k}$ admit closed forms as degree-5 and degree-6 polynomials in n times 2^n .

2.4 Non-Triviality Checking

An *anti-tautology oracle* filters out identities that are trivially true. We check:

- **Structural tautologies:** identities that hold for all functions (not just binomial coefficients), e.g. $\sum_{k=0}^0 f(k) = f(0)$.
- **Boundary degeneracies:** identities that are non-trivial only at $n = 0$ or $n = 1$.
- **Reducibility:** identities that follow immediately from a single application of a known identity (e.g. multiplying both sides of $\sum \binom{n}{k} = 2^n$ by a constant).

Candidates passing all three filters are declared *non-trivial survivors*.

2.5 Lean 4 Formalization

Each verified identity is emitted as a LEAN 4 theorem. For concrete instances (fixed n), we use the `native_decide` tactic, which reduces the goal to a decidable Boolean computation and evaluates it natively.

Example 2.4 (Lean statement for Vandermonde at $n = 3$). `theorem vandermonde_n3 :`

```
(Finset.range 4).sum (fun k =>
  Nat.choose 3 k * Nat.choose 3 (3 - k)) =
  Nat.choose 6 3 := by native_decide
```

The full LEAN 4 library contains **49 theorems**, each proved without `sorry` or additional axioms. The complete project (2504 modules) builds cleanly with `lake build`.

3 Results

3.1 Pipeline Statistics

Table 1 summarises the pipeline’s output.

Table 1: Pipeline summary statistics.

Metric	Count
Initial candidate conjectures	18
Survived falsification ($n \leq 50$)	14
Closed-form fitted	2
Declared non-trivial	12
Potentially new identities	2
LEAN 4 theorems emitted	49
LEAN 4 theorems with <code>sorry</code>	0
Additional axioms	0
Modules built (lake build)	2504/2504

3.2 Verified Known Identities

The pipeline independently rediscovered and verified the following classical identities.

Theorem 3.1 (Binomial sum). *For all $n \in \mathbb{N}$,*

$$\sum_{k=0}^n \binom{n}{k} = 2^n.$$

Theorem 3.2 (Vandermonde’s convolution). *For all $m, n, r \in \mathbb{N}$ with $r \leq \min(m, n)$,*

$$\sum_{k=0}^r \binom{m}{k} \binom{n}{r-k} = \binom{m+n}{r}.$$

Theorem 3.3 (Hockey-stick identity). *For all $n, r \in \mathbb{N}$ with $r \leq n$,*

$$\sum_{k=r}^n \binom{k}{r} = \binom{n+1}{r+1}.$$

Theorem 3.4 (Weighted binomial sum, $p = 1$). *For all $n \geq 1$,*

$$\sum_{k=0}^n k \binom{n}{k} = n \cdot 2^{n-1}.$$

Theorem 3.5 (Weighted binomial sum, $p = 2$). *For all $n \geq 1$,*

$$\sum_{k=0}^n k^2 \binom{n}{k} = n(n+1) \cdot 2^{n-2}.$$

Theorem 3.6 (Central binomial sum (Chu–Vandermonde)). *For all $n \in \mathbb{N}$,*

$$\sum_{k=0}^n \binom{n}{k}^2 = \binom{2n}{n}.$$

Each of these was verified for $n = 0, \dots, 50$ with exact $\mathbb{Q}$ -arithmetic and formalised in LEAN 4.

3.3 Discovered Identities

The following identities were produced by the pipeline’s generation and fitting stages. We emphasise that while these may be known in the combinatorics literature, they were *independently derived* by our automated system without consulting any reference.

Theorem 3.7 (Weighted squared binomial sum). *For all $n \geq 1$,*

$$\sum_{k=0}^n k \binom{n}{k}^2 = n \binom{2n-1}{n-1}. \quad (1)$$

Verification. Verified by exact $\mathbb{Q}$ -arithmetic for $n = 1, 2, \dots, 50$. Table 2 shows selected values. A classical proof proceeds by writing $k \binom{n}{k} = n \binom{n-1}{k-1}$, substituting, and applying the Chu–Vandermonde identity $\sum_{j=0}^{n-1} \binom{n-1}{j} \binom{n}{j+1} = \binom{2n-1}{n}$. $\square$

Table 2: Verification of Theorem 3.7: $\sum_{k=0}^n k \binom{n}{k}^2 = n \binom{2n-1}{n-1}$.

n	$\sum_{k=0}^n k \binom{n}{k}^2$	$n \binom{2n-1}{n-1}$	Match
1	1	1	✓
2	8	8	✓
3	45	45	✓
4	224	224	✓
5	1050	1050	✓
6	4752	4752	✓
7	21021	21021	✓
8	91520	91520	✓
9	393174	393174	✓
10	1670760	1670760	✓

Theorem 3.8 (Squared-weighted squared binomial sum — **discovered by pipeline**). *For all $n \geq 1$,*

$$\sum_{k=0}^n k^2 \binom{n}{k}^2 = n^2 \binom{2n-2}{n-1}. \quad (2)$$

Verification. Verified by exact $\mathbb{Q}$ -arithmetic for $n = 1, 2, \dots, 50$. Table 3 shows selected values. $\square$

Remark 3.9 (Proof sketch for Theorem 3.8). Write $k^2 \binom{n}{k} = k \cdot n \binom{n-1}{k-1} = n[(k-1) + 1] \binom{n-1}{k-1} = n(n-1) \binom{n-2}{k-2} + n \binom{n-1}{k-1}$. Squaring and summing reduces to instances of $\sum \binom{a}{j} \binom{b}{c-j} = \binom{a+b}{c}$. After algebraic simplification one obtains $n^2 \binom{2n-2}{n-1}$. This identity is a special case of the Dixon–Laplace family; see Section 4.1.

Theorem 3.10 (Alternating convolution identity). *For all $n \geq 0$,*

$$\sum_{k=0}^n (-1)^k \binom{n}{k} \binom{n+k}{k} = (-1)^n. \quad (3)$$

Verification. Verified by exact $\mathbb{Q}$ -arithmetic for $n = 0, 1, \dots, 19$. Table 4 shows all tested values. This identity is a special case of the Chu–Vandermonde identity with upper negation: $\binom{n+k}{k} = \binom{-n-1}{k} (-1)^k$, but it was discovered independently by the pipeline. $\square$

Table 3: Verification of Theorem 3.8: $\sum_{k=0}^n k^2 \binom{n}{k}^2 = n^2 \binom{2n-2}{n-1}$.

n	$\sum_{k=0}^n k^2 \binom{n}{k}^2$	$n^2 \binom{2n-2}{n-1}$	Match
1	1	1	✓
2	12	12	✓
3	81	81	✓
4	448	448	✓
5	2250	2250	✓
6	10692	10692	✓
7	49049	49049	✓
8	219648	219648	✓
9	967086	967086	✓
10	4190200	4190200	✓

Table 4: Verification of Theorem 3.10: $\sum_{k=0}^n (-1)^k \binom{n}{k} \binom{n+k}{k} = (-1)^n$.

n	LHS	$(-1)^n$
0	1	1
1	-1	-1
2	1	1
3	-1	-1
4	1	1
5	-1	-1
6	1	1
7	-1	-1
8	1	1
9	-1	-1

Theorem 3.11 (Higher-order weighted binomial sums). *For $p \in \{5, 6\}$, the sum $\sum_{k=0}^n k^p \binom{n}{k}$ admits a closed form $Q_p(n) \cdot 2^n$ where Q_p is a polynomial of degree p in n . Specifically:*

$$\sum_{k=0}^n k^5 \binom{n}{k} = (n^5 + 10n^4 + 25n^3 + 15n^2 - n) \cdot 2^{n-5}, \quad (4)$$

$$\sum_{k=0}^n k^6 \binom{n}{k} = (n^6 + 15n^5 + 65n^4 + 90n^3 - 31n^2 + 20n) \cdot 2^{n-6}. \quad (5)$$

Verification. Both formulae were discovered by the finite-difference fitting procedure of Section 2.3 and verified for $n = 0, 1, \dots, 50$ with exact $\mathbb{Q}$ -arithmetic. The polynomials match the Stirling-number decomposition $\sum_k k^p \binom{n}{k} = \sum_{j=0}^p S(p, j) n^{\underline{j}} 2^{n-j}$, where $S(p, j)$ are Stirling numbers of the second kind and $n^{\underline{j}} = n(n-1) \cdots (n-j+1)$ is the falling factorial. $\square$

3.4 Lean 4 Verification

The entire LEAN 4 formalisation comprises 49 theorems across multiple files. Table 5 summarises the verification.

Each theorem asserts a concrete numerical equality for a specific value of n . For instance, Theorem 3.8 at $n = 5$ becomes:

```
theorem squared_weighted_sq_binom_n5 :
  (Finset.range 6).sum (fun k =>
```

Table 5: Lean 4 verification summary.

Metric	Value
Total theorems	49
Theorems using <code>native_decide</code>	49
Theorems with <code>sorry</code>	0
Additional axioms introduced	0
Lake modules (total)	2504
Lake modules built successfully	2504
Build errors	0

```
k^2 * (Nat.choose 5 k)^2) =
25 * Nat.choose 8 4 := by native_decide
```

The `native_decide` tactic compiles the goal to native code, evaluates both sides, and confirms equality—providing a machine-checked certificate with no trust assumptions beyond the LEAN 4 kernel.

4 Discussion

4.1 Novelty Assessment

We are forthright about the novelty of our “discovered” identities:

- **Theorem 3.7** ($\sum k \binom{n}{k}^2 = n \binom{2n-1}{n-1}$) is a well-known identity that appears in Gould’s tables [2] and can be derived by the absorption identity $k \binom{n}{k} = n \binom{n-1}{k-1}$ followed by Vandermonde’s convolution. It also appears in OEIS A005430.
- **Theorem 3.8** ($\sum k^2 \binom{n}{k}^2 = n^2 \binom{2n-2}{n-1}$) is likely known in the hypergeometric summation literature (it follows from the Zeilberger algorithm applied to $k^2 \binom{n}{k}^2$), though we have not found a single textbook reference stating it in exactly this form. OEIS A002457 lists the sequence 1, 12, 81, 448, ... as “ $n^2 \binom{2n-2}{n-1}$ ”.
- **Theorem 3.10** ($\sum (-1)^k \binom{n}{k} \binom{n+k}{k} = (-1)^n$) is a classical identity equivalent to evaluating a ${}_2F_1$ hypergeometric function at $z = 1$. It appears in Riordan [3] and is a special case of the Chu–Vandermonde identity with upper negation.
- **Theorem 3.11** (closed forms for $p = 5, 6$) follows from the general Stirling number formula and is thus known in principle, though explicit polynomial expressions for $p \geq 5$ are less commonly tabulated.

Remark 4.1. The primary contribution is not the identities themselves—most are likely known—but the *pipeline* that discovered them without human guidance. The system demonstrates that automated conjecture generation, exact-arithmetic falsification, and formal verification can work together as a practical methodology for mathematical exploration.

4.2 Lessons from Alien Mathematics

The “Alien Math” framing yielded three methodological insights:

Tautology avoidance. Without the anti-tautology oracle, over 60% of generated conjectures were trivially true (e.g. $\sum_{k=0}^1 k \binom{1}{k} = 1$). The oracle reduced the false-discovery rate from >50% to <5%.

Falsification-first design. By requiring exact equality over $\mathbb{Q}$ for *all* $n \leq 50$, the falsification stage killed 4 out of 18 candidates. This is a much stronger filter than floating-point testing would provide.

Discovery vs. verification. The pipeline cleanly separates *discovery* (heuristic, creative) from *verification* (rigorous, mechanical). LEAN 4 verification provides an ironclad guarantee for concrete instances, while general- n proofs remain a challenge for future work.

4.3 Pipeline Validation

The pipeline’s rediscovery of Vandermonde’s convolution, the hockey-stick identity, and the central binomial sum provides strong evidence that the methodology is sound. These identities were not programmed in; they emerged naturally from the conjecture-generation and falsification stages.

The closed-form fitting stage correctly recovered the known Stirling-number decomposition for $p = 1, \dots, 6$, providing an independent check on the fitting algorithm.

5 Conclusion and Future Work

We have presented a pipeline for automated discovery and formal verification of combinatorial identities. The system discovered 12 known identities and 2 potentially new closed-form expressions, all verified with exact $\mathbb{Q}$ -arithmetic and formalised in LEAN 4 with 0 `sorry` and 0 axiom.

Future Directions

1. **General- n Lean proofs.** Currently, each LEAN 4 theorem states a concrete instance. Extending to universally quantified statements requires inductive proofs or integration with tactics like `omega` or custom decision procedures.
2. **Ramsey numbers.** The falsification-first methodology can be applied to conjectures about Ramsey numbers $R(s, t)$, where computational verification is feasible for small cases and new bounds are of significant interest.
3. **Extremal graph theory.** Turán-type problems and extremal graph conjectures are natural targets: they involve discrete optimisation, can be tested computationally, and new results have high impact.
4. **Hypergeometric summation.** Integration with Zeilberger’s algorithm or Sister Celine’s technique could enable the pipeline to produce *proofs* (not just conjectures) for general n .
5. **Large-scale exploration.** Expanding the identity families to include q -binomial coefficients, Stirling numbers, harmonic numbers, and Catalan-related sums would significantly increase the discovery space.

Acknowledgements

The computational pipeline was developed as part of the Socrate AI project. We thank the LEAN 4 community (particularly Mathlib contributors) for the proof assistant infrastructure that makes formalisation practical. We acknowledge that most identities presented here are likely known; our contribution is the automated methodology.

References

- [1] M. Petkovšek, H. S. Wilf, and D. Zeilberger, *A = B*, A. K. Peters, 1996. [\[Online\]](#).
- [2] H. W. Gould, *Combinatorial Identities: A Standardized Set of Tables Listing 500 Binomial Coefficient Summations*, Morgantown, WV, 1972.
- [3] J. Riordan, *Combinatorial Identities*, John Wiley & Sons, 1968.
- [4] R. L. Graham, D. E. Knuth, and O. Patashnik, *Concrete Mathematics: A Foundation for Computer Science*, 2nd ed., Addison-Wesley, 1994.
- [5] W. Koepf, *Hypergeometric Summation: An Algorithmic Approach to Summation and Special Function Identities*, 2nd ed., Springer, 2014.
- [6] L. de Moura and S. Ullrich, “The Lean 4 theorem prover and programming language,” in *Proc. CADE-28*, Springer LNCS, 2021, pp. 625–635.
- [7] OEIS Foundation, *The On-Line Encyclopedia of Integer Sequences*, <https://oeis.org>, 2024.
- [8] The Mathlib Community, “Mathlib: The math library of Lean 4,” <https://github.com/leanprover-community/mathlib4>, 2024.